

What is a MapReduce?

Michael Kleber

with many slides shamelessly stolen from Jeff Dean and Yonatan Zunger



Word Count and friends

Input: (URL, contents)

Map output

("the", "1")

("the", "www.bar.org/index.html")

(hostname, doc term frequencies)

("purple", "cow")

Reduce output

("the", "1048576")

("the", list of URLs)

(hostname, site term frequencies)

("purple", probability distrib of
next word in English)

Mappers can take command-line arguments, eg a regular expression.


("line number", "matching line")

("line number", "matching line")


With 1800 machines, MR_Grep scanned 1 terabyte in 100 seconds.

Query Frequency Over Time

Queries containing “eclipse”




Queries containing “world series”




Queries containing “full moon”



Queries containing “summer olympics”



Queries containing “watermelon”



Queries containing “Opteron”



Sorting

In addition to map and reduce functions, you may specify

Partition: (k' , number of reducers) \mapsto choice of reducer for k'

Default partition is $(\text{hash}(k') \bmod \# \text{reducers})$, for load balancing, but:

- Output file for k' reduction determined by its partition
- Guarantee: each reducer sees the keys in its partition in sorted order (implemented by the invisible shuffle-and-sort stage)

Map: produces (sort key, record)

Partition: send consecutive blocks of sort keys to same reducer, e.g. by using most significant bits of sort keys

Reduce: identity function

MR_Sort sorted 1 terabyte of 100-byte records in 14 minutes

Structure of the Web

Input is (URL, contents) again

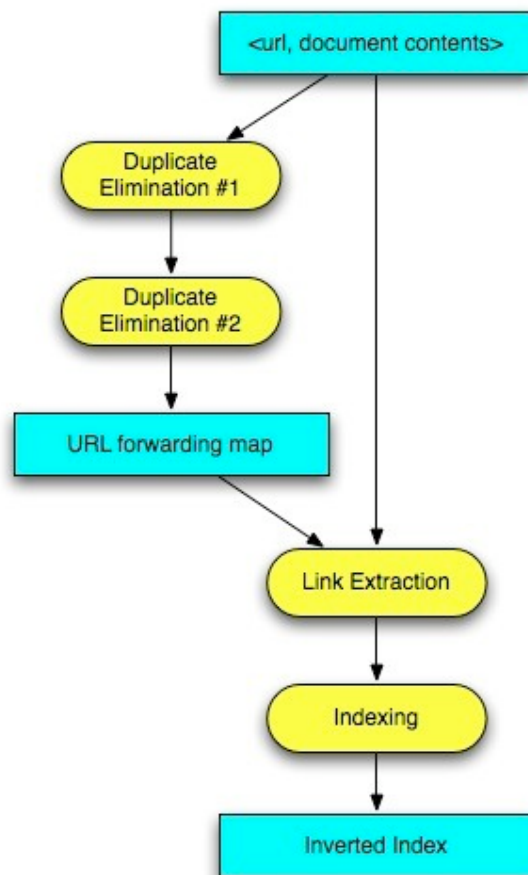
Scan through the document's contents looking for links to other URLs

- Map outputs (URL, linked-to URL)
you get a simple representation of the WWW “link graph”
- Map outputs (linked-to URL, URL)
you get the reverse link graph, “what web pages link to me?”
- Map outputs (linked-to URL, anchor text)
you get “how do other web pages characterize me?”

Google uses this “anchor text propagation” in indexing.



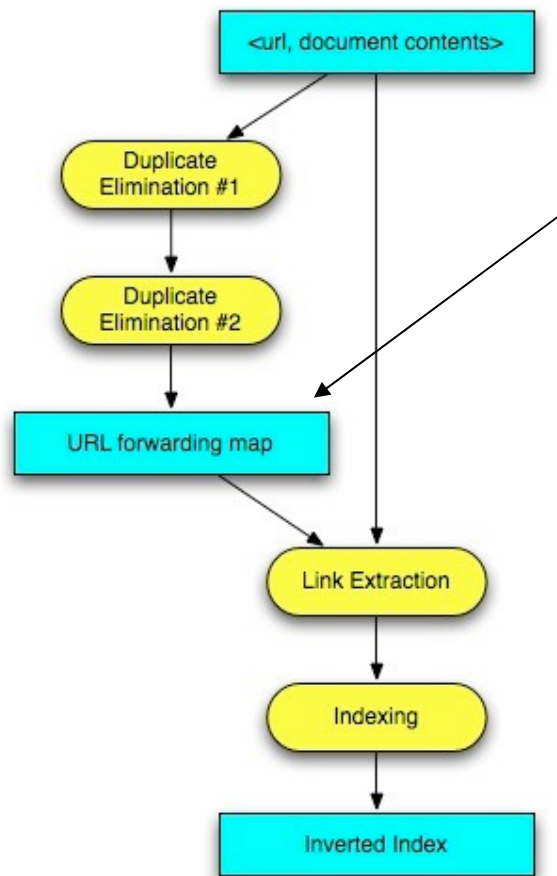
Example: A Simple Indexing Pipeline



Input: (url, document) pairs

Output: Inverted index including anchor text

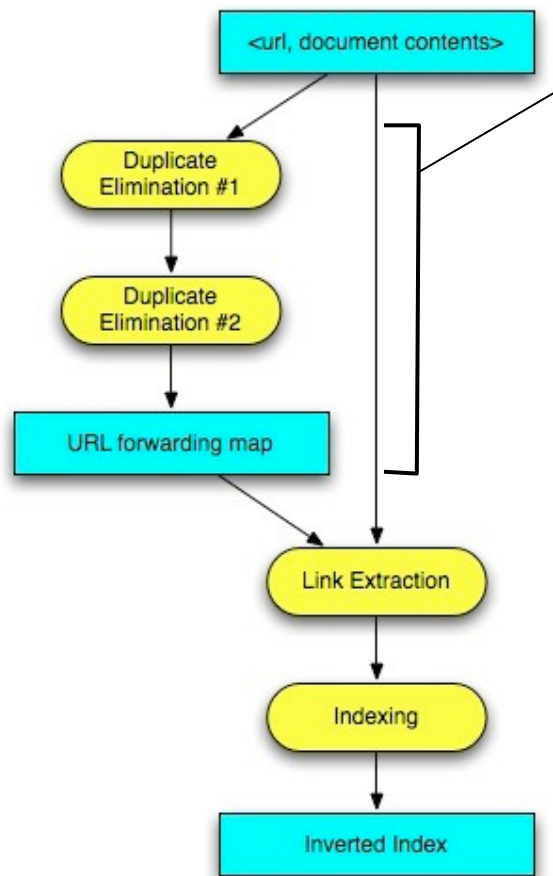
Example: A Simple Indexing Pipeline



Duplicate Elimination

Want: a sorted table of
(original URL, canonical URL)
for each URL \neq its canonical URL

Example: A Simple Indexing Pipeline

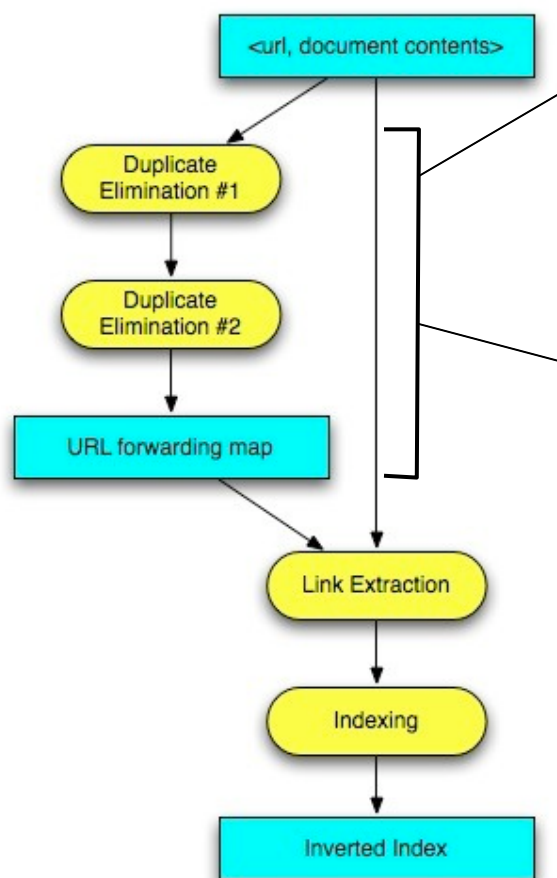


Duplicate Elimination

Map: (url, contents) \mapsto <(contentfp, url + aux info)>

Reduce: (contentfp, <url + info>) \mapsto
<(contentfp, url + "best" url)>

Example: A Simple Indexing Pipeline



Duplicate Elimination

Map: (url, contents) \mapsto <(contentfp, url + aux info)>

Reduce: (contentfp, <url + info>) \mapsto
<(contentfp, url + "best" url)>

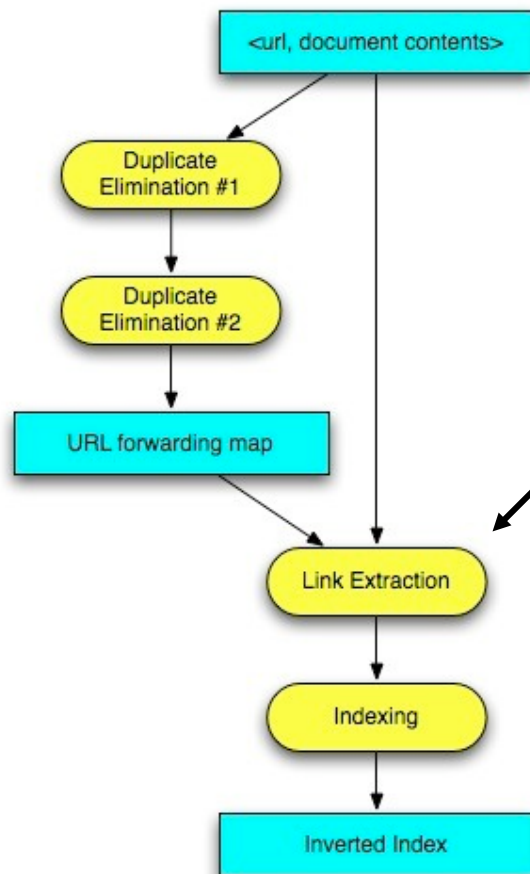
Reorganize that

Map: (contentfp, url + canonical url) \mapsto
(url, canonical url) (or nothing if equal)

Reduce: Identity

Partition: hash(url) % number of output shards

Example: A Simple Indexing Pipeline



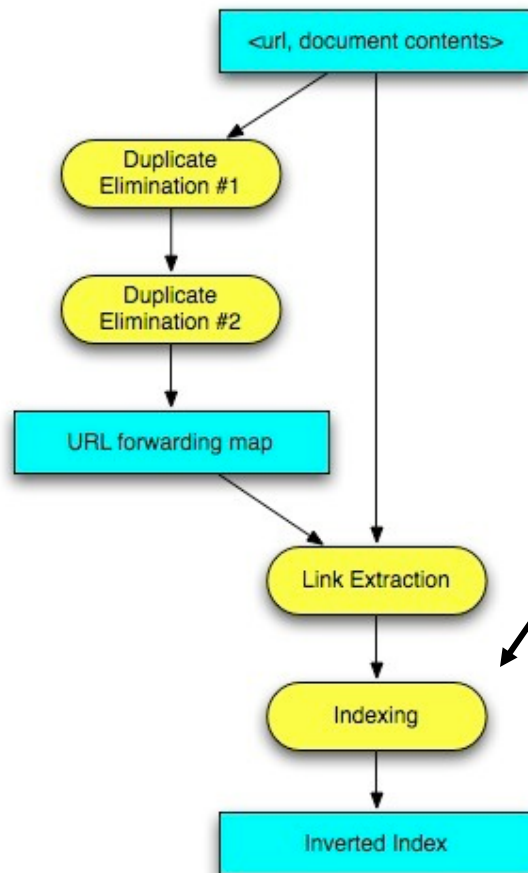
Link Extraction

Map: (url, contents)

- (a) Parse contents, forward each link using map; output (target url, anchor info)
- (b) Try to forward doc URL using map; if it doesn't forward, output (url, contents)

Reduce: (forwarded url, <content or anchor>) \mapsto (url, contents + anchors)

Example: A Simple Indexing Pipeline



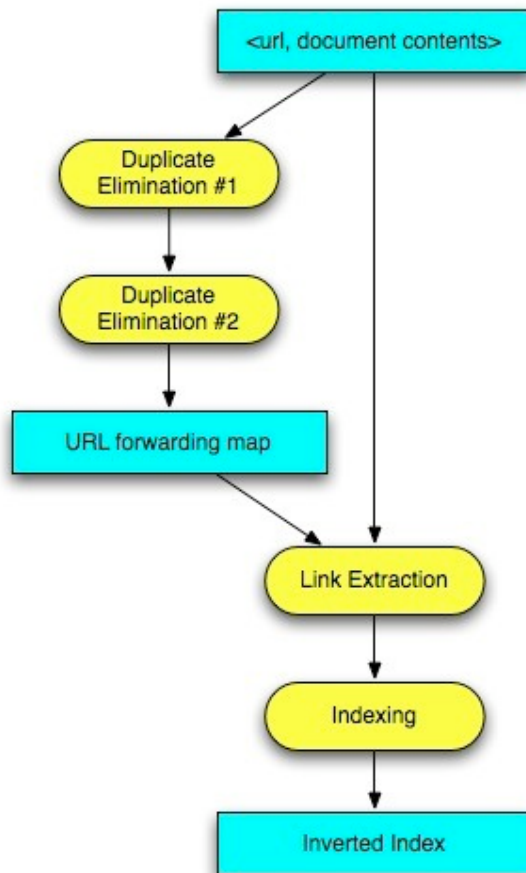
Indexing

Map: (url, merged contents) \mapsto \langle (term, occurrence info) \rangle

Partition: $\text{hash}(\text{url}) \% \text{number of output shards}$

Reduce: (term, occurrences in partition of docs) \mapsto (term, compressed posting list)

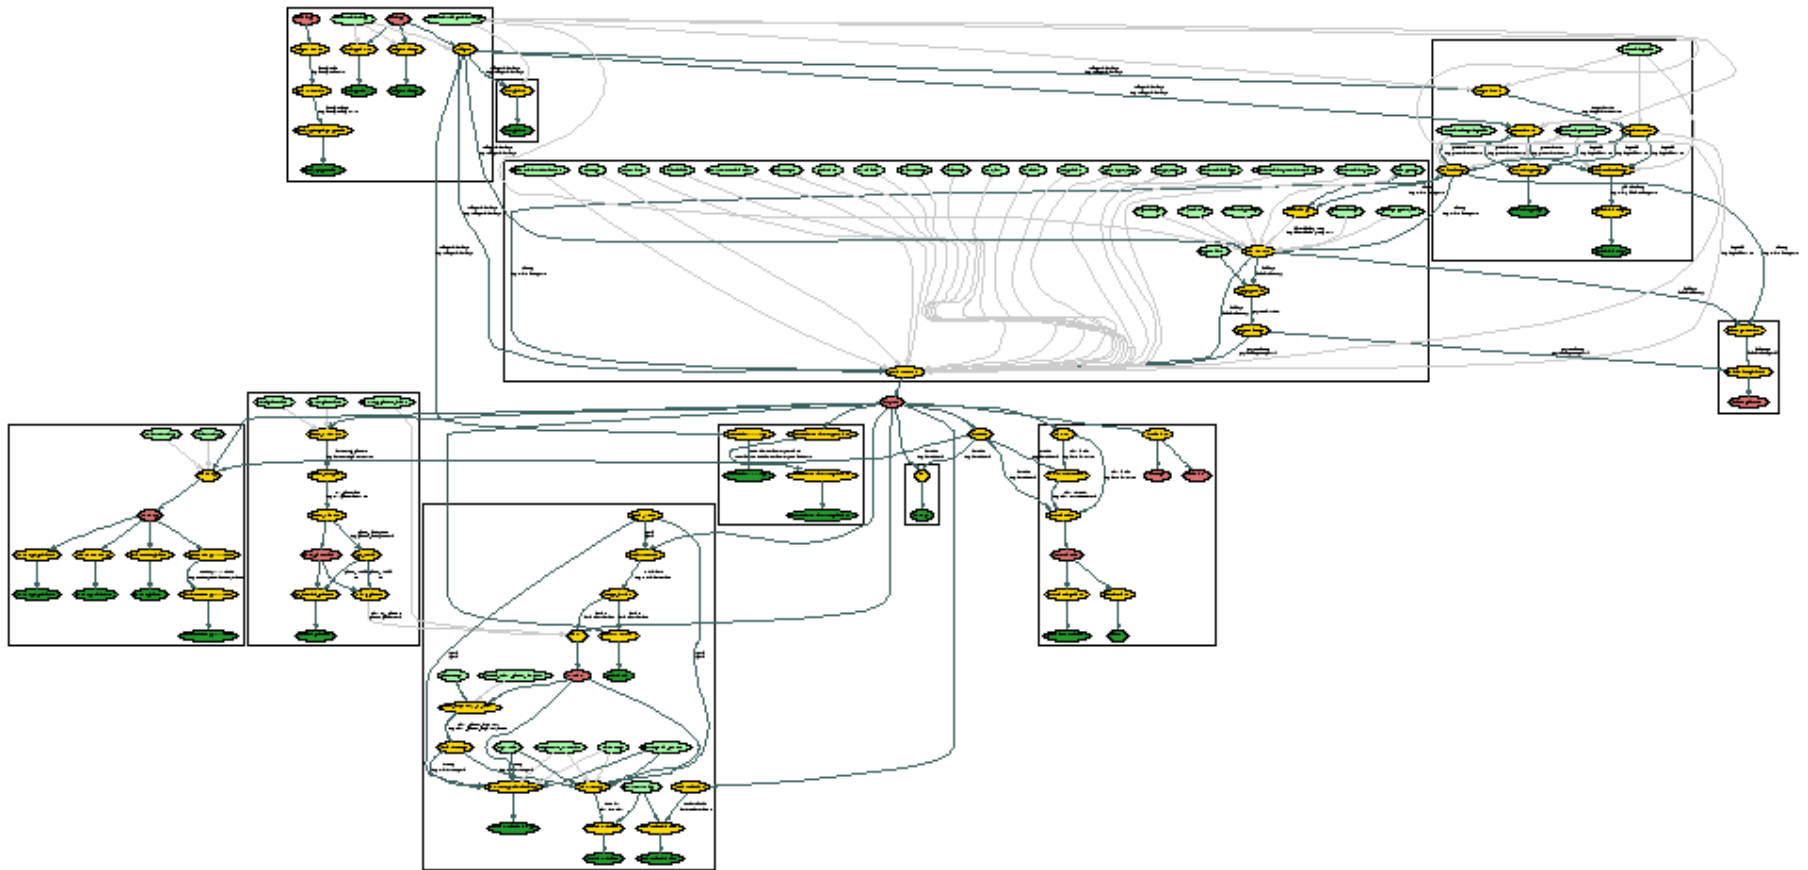
Example: A Simple Indexing Pipeline



Final product: a map from each term to its compressed posting list

Each partition file contains docs with fixed $(\text{hash}(\text{url}) \% \text{num files})$

Real indexing pipelines are a bit messier...



PageRank

Indexing pipeline can answer the question

What are all the pages that match this query?

But most people don't want to look at all million hits. We need to answer

What are the best pages that match this query?

PageRank: use the link structure of the web to find “best” pages



PageRank

Suppose you browsed the web randomly, i.e. by clicking on random links.
What fraction of your time would you spend on each page?


$T(X)$ = time spent on page X

$L(X)$ = # links from page X to other pages


Then the Ts and Ls would satisfy equations like

$$T(X) = T(A)/L(A) + T(B)/L(B) + T(C)/L(C)$$

if A, B, C are all the pages linking to X



PageRank



$$T(X) = T(A)/L(A) + T(B)/L(B) + T(C)/L(C)$$

Problems:

- Pages with no outbound links = black holes
- How to calculate? Is there even a unique solution?

Solution to both: damping factor ($d = .85$)

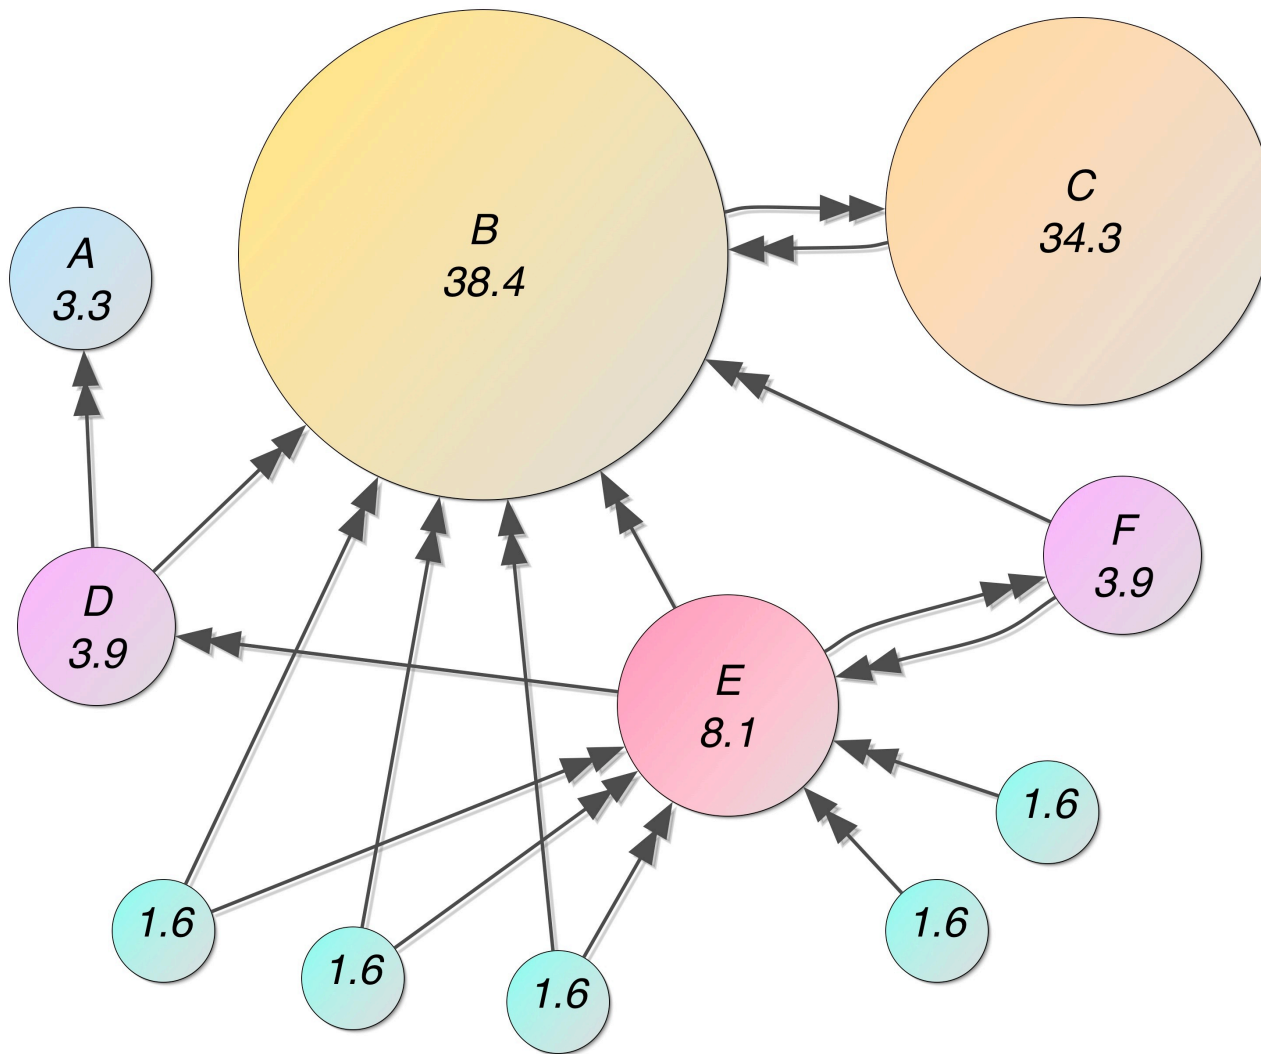
$$PR(X) = (1-d) + d*(PR(A)/L(A) + PR(B)/L(B) + PR(C)/L(C))$$

Random browsing, except with probability $1-d$ you jump to a random page.

Brin & Page, "The Anatomy of a Large-Scale Hypertextual Web Search Engine" (1998)




PageRank



PageRank

$$PR(X) = (1-d) + d * (PR(A)/L(A) + PR(B)/L(B) + PR(C)/L(C))$$



We can calculate PageRank iteratively.

$PR_i(X)$ = probability of being on page X after i random steps

- Initial state:

$$PR_0(X) = 1 \text{ for all pages } X.$$

- Update:

$$PR_{i+1}(X) = (1-d) + d \sum_{Y \rightarrow X} PR_i(Y)/L(Y)$$

- $PR_i(X)$ converges to $PR(X)$ as $i \rightarrow \infty$. Keep iterating.

PageRank: MapReduce Implementation

$$PR(X) = (1-d) + d*(PR(A)/L(A) + PR(B)/L(B) + PR(C)/L(C))$$

- Perform one MR to produce the link graph and set up initial state

$$(\text{URL}, \text{contents}) \mapsto (\text{URL}, [1.0, \text{list of linked-to URLs}])$$

- Each MapReduce performs one step of the iteration
The mapper and reducer keys are both the set of all URLs

$$\begin{aligned} \text{Map: } (Y, [PR(Y), Z_1 \dots Z_n]) &\mapsto (Z_i, PR(Y)/n) \quad \text{for } i=1, \dots, n \\ & (Y, [Z_1 \dots Z_n]) \quad \text{remember the link graph} \end{aligned}$$

$$\text{Reduce: } (Y, \langle S_0, S_1, \dots, S_m, [Z_1 \dots Z_n] \rangle) \mapsto (Y, [PR(Y), Z_1 \dots Z_n])$$

- An outside controller runs each MR iteration and decides when to stop.